

DATA COLLECTION

```
In [1]: # import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # To Import Dataset
sd=pd.read_csv(r"c:\Users\user\Downloads\placement.csv")
sd
```

```
Out[2]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
...
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

1000 rows × 3 columns

```
In [3]: # to display top 10 rows
sd.head(10)
```

```
Out[3]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
5	7.30	23.0	1
6	6.69	11.0	0
7	7.12	39.0	1
8	6.45	38.0	0
9	7.75	94.0	1

DATA CLEANING AND PRE_PROCESSING

```
In [4]: sd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cgpa                  1000 non-null  float64
1   placement_exam_marks 1000 non-null  float64
2   placed                1000 non-null  int64
dtypes: float64(2), int64(1)
memory usage: 23.6 KB
```

```
In [5]: # to display summary of statistics
sd.describe()
```

```
Out[5]:
```

	cgpa	placement_exam_marks	placed
count	1000.000000	1000.000000	1000.000000
mean	6.961240	32.225000	0.489000
std	0.615898	19.130822	0.500129
min	4.890000	0.000000	0.000000
25%	6.550000	17.000000	0.000000
50%	6.960000	28.000000	0.000000
75%	7.370000	44.000000	1.000000
max	9.120000	100.000000	1.000000

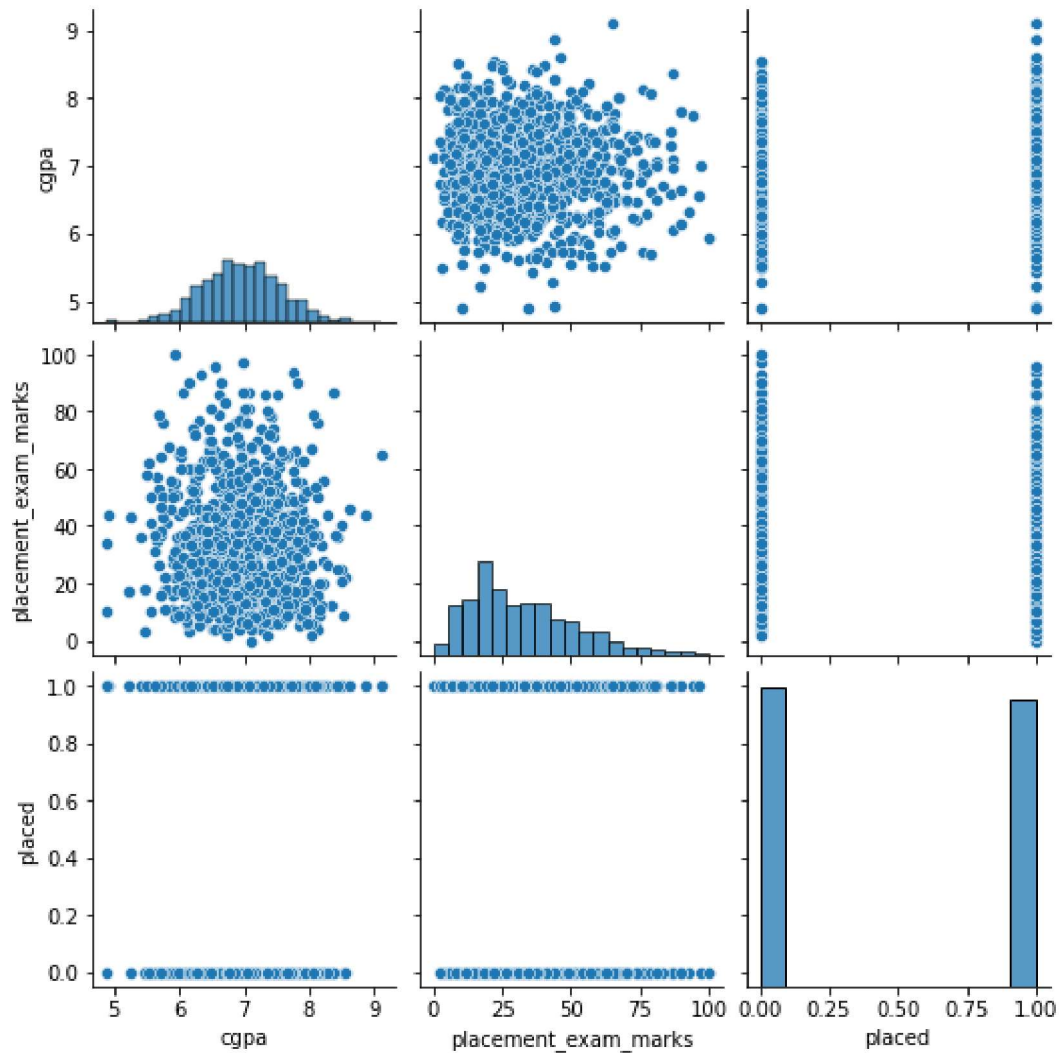
```
In [6]: #to display colums heading  
sd.columns
```

```
Out[6]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

EDA and visualization

```
In [7]: sns.pairplot(sd)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x220767c6fd0>
```




```
In [8]: sns.distplot(sd['cgpa'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

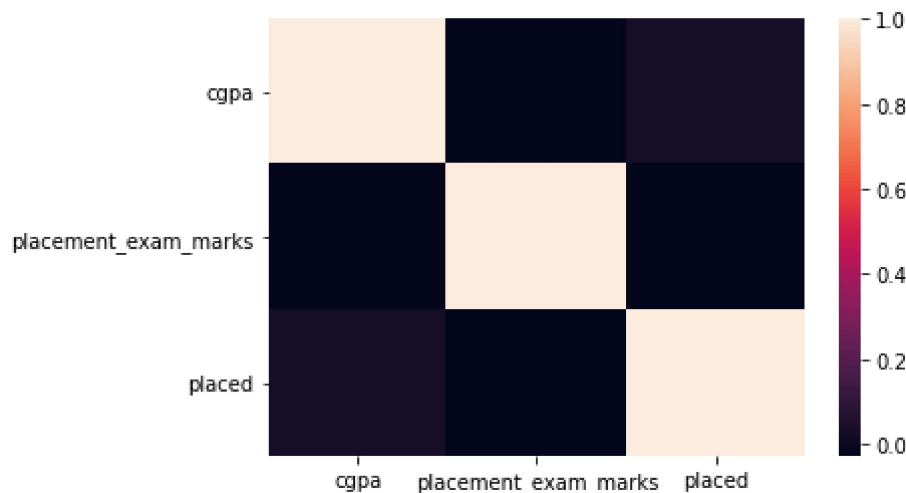
```
Out[8]: <AxesSubplot:xlabel='cgpa', ylabel='Density'>
```



```
In [9]: sd1=sd[['cgpa', 'placement_exam_marks', 'placed']]
```

```
In [10]: sns.heatmap(sd1.corr())
```

```
Out[10]: <AxesSubplot:>
```



TO TRAIN THE MODEL _MODEL BUILDING

we are going to train a Linear Regression model; we need to split out the data into two variables x and y where x is independent on x (output) and y is dependent on x (output) address column as it is not required our model

```
In [11]: x= sd1[['cgpa', 'placement_exam_marks']]
y=sd1['placed']
```

```
In [12]: # To split my dataset into training data and test data
from sklearn .model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4)
```

```
In [13]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[13]: LinearRegression()

```
In [14]: print(lr.intercept_)

0.5099041525907622
```


```
In [15]: coeff= pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[15]:

	Co-efficient
cgpa	0.000756
placement_exam_marks	-0.001088

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x22078f80520>



```
In [17]: print(lr.score(x_test,y_test))

-0.003601326842912256
```

```
In [18]: lr.score(x_train,y_train)
```

```
Out[18]: 0.0017995666628813911
```

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: dr=Ridge(alpha=10)  
dr.fit(x_train,y_train)
```

```
Out[20]: Ridge(alpha=10)
```

```
In [21]: dr.score(x_test,y_test)
```

```
Out[21]: -0.003605905461431025
```

```
In [22]: dr.score(x_train,y_train)
```

```
Out[22]: 0.0017995651466249374
```

```
In [23]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[23]: Lasso(alpha=10)
```

```
In [24]: la.score(x_test,y_test)
```

```
Out[24]: -0.0020250506262655676
```

```
In [25]: la.score(x_train,y_train)
```

```
Out[25]: 0.0
```

ElasticNet

```
In [26]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[26]: ElasticNet()
```

```
In [27]: print(en.coef_)
```

```
[ 0. -0.]
```

```
In [28]: print(en.intercept_)
```

```
0.48
```

```
In [29]: prediction=en.predict(x_test)
```

```
In [30]: print(en.score(x_test,y_test))
```

```
-0.0020250506262655676
```

Evaluation metrics

```
In [31]: from sklearn import metrics
```

```
In [33]: print("mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
mean Absolute Error: 0.5001000000000001
```

```
In [34]: print("mean squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
mean squared Error: 0.2505
```

```
In [35]: print("Root mean Absolytre Error:",np.sqrt(metrics.mean_squared_error(y_test,pr
```

```
Root mean Absolytre Error: 0.500499750249688
```

```
In [ ]:
```